

```

In [1]: import numpy as np

def relu(x):
    return np.maximum(0, x)

def relu_deriv(x):
    return (x > 0).astype(float)

def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def sigmoid_deriv(x):
    s = sigmoid(x)
    return s * (1 - s)

def forward(x, parameters, activation='relu'):
    caches = []
    A = x
    L = len(parameters) // 2

    for l in range(1, L + 1):
        W = parameters["W" + str(l)]
        b = parameters["b" + str(l)]
        Z = np.dot(W, A) + b
        caches.append((A, Z))
        if activation == 'relu':
            A = relu(Z)
        elif activation == 'sigmoid':
            A = sigmoid(Z)
        else:
            raise ValueError("Unsupported activation function")

    AL = A
    return AL, caches

def backward(dAL, caches, parameters, activation='relu'):

    grads = {}
    L = len(caches)

    A_prev, Z = caches[-1]
    if activation == 'relu':
        dZ = dAL * relu_deriv(Z)
    elif activation == 'sigmoid':
        dZ = dAL * sigmoid_deriv(Z)
    else:
        raise ValueError("Unsupported activation function")

    grads["dW" + str(L)] = np.einsum('ik,jk->ij', dZ, A_prev)
    grads["db" + str(L)] = np.sum(dZ, axis=1, keepdims=True)
    dA_prev = np.dot(parameters["W" + str(L)].T, dZ)

    for l in range(L - 1, 0, -1):
        A_prev, Z = caches[l - 1]
        if activation == 'relu':
            dZ = dA_prev * relu_deriv(Z)
        elif activation == 'sigmoid':
            dZ = dA_prev * sigmoid_deriv(Z)
        grads["dW" + str(l)] = np.einsum('ik,jk->ij', dZ, A_prev)
        grads["db" + str(l)] = np.sum(dZ, axis=1, keepdims=True)
        if l > 1:
            dA_prev = np.dot(parameters["W" + str(l)].T, dZ)

```

```
return grads
```

A): Warm up:

0
0 0
0 0 0
0 0 0 0
0 0 0
0

L_1 L_2 L_3 L_4

①: from raw data to L_1 :
 $L_1 = W_{0 \rightarrow 1} \cdot L_0 + \delta_{0 \rightarrow 1}$
 6×6 6

②
 $L_2 = W_{1 \rightarrow 2} \cdot L_1 + \delta_{1 \rightarrow 2}$
Dim: 4×6 4

③ $L_3 = W_{2 \rightarrow 3} \cdot L_2 + \delta_{2 \rightarrow 3}$
Dim: 3×4 3

④ $L_4 = W_{3 \rightarrow 4} \cdot L_3 + \delta_{3 \rightarrow 4}$
 2×3 2

In [2]: # --- Example Usage ---

```
# Define network architecture: input 6, then layers with 6, 4, 3, and 2 neurons.
layer_dims = [6, 6, 4, 3, 2]
parameters = {}
np.random.seed(42)
L = len(layer_dims) - 1

for l in range(1, L + 1):
    parameters["W" + str(l)] = np.random.randn(layer_dims[l], layer_dims[l - 1]) * 0.01
    parameters["b" + str(l)] = np.zeros((layer_dims[l], 1))

x = np.random.randn(6, 1)

# Forward pass
AL, caches = forward(x, parameters, activation='sigmoid')
print("Forward propagation output (AL):\n", AL)

# Assume a dummy derivative from the loss (e.g., from mean squared error)
# For our output layer with 2 neurons, dAL has shape (2, 1)
dAL = np.random.randn(2, 1)

# Backward pass to compute gradients
grads = backward(dAL, caches, parameters, activation='sigmoid')

# Print gradients for each layer
for l in range(1, L + 1):
    print(f"\nGradients for Layer {l}:")
    print("dW" + str(l) + ":\n", grads["dW" + str(l)])
    print("db" + str(l) + ":\n", grads["db" + str(l)])
```

Forward propagation output (AL):

```

[[0.49863851]
 [0.50075413]]

Gradients for Layer 1:
dW1:
[[ 1.26488930e-09 -2.73979226e-08 -3.02809797e-09  4.92266836e-09
  2.03722379e-08 -7.14416856e-09]
 [ 4.16995430e-10 -9.03225959e-09 -9.98271557e-10  1.62285364e-09
  6.71610557e-09 -2.35521451e-09]
 [ 1.06951551e-09 -2.31660615e-08 -2.56038038e-09  4.16231695e-09
  1.72255583e-08 -6.04068600e-09]
 [ 3.36203080e-10 -7.28227044e-09 -8.04857674e-10  1.30842775e-09
  5.41486840e-09 -1.89889460e-09]
 [ 4.95193484e-10 -1.07260554e-08 -1.18547479e-09  1.92718311e-09
  7.97555915e-09 -2.79688168e-09]
 [ 1.07490250e-11 -2.32827455e-10 -2.57327664e-11  4.18328192e-11
  1.73123209e-10 -6.07111204e-11]]
db1:
[[1.37846403e-08]
 [4.54437555e-09]
 [1.16554758e-08]
 [3.66390840e-09]
 [5.39657032e-09]
 [1.17141827e-10]]

Gradients for Layer 2:
dW2:
[[-4.64793566e-08 -4.60100809e-08 -4.69072904e-08 -4.70964132e-08
 -4.62203650e-08 -4.58394542e-08]
 [-2.48009333e-06 -2.45505324e-06 -2.50292747e-06 -2.51301888e-06
 -2.46627379e-06 -2.44594877e-06]
 [ 1.04794593e-06  1.03736541e-06  1.05759433e-06  1.06185838e-06
  1.04210658e-06  1.03351838e-06]
 [ 2.06265479e-06  2.04182933e-06  2.08164558e-06  2.09003844e-06
  2.05116129e-06  2.03425728e-06]]
db2:
[[-9.27281473e-08]
 [-4.94788390e-06]
 [ 2.09069341e-06]
 [ 4.11507755e-06]]

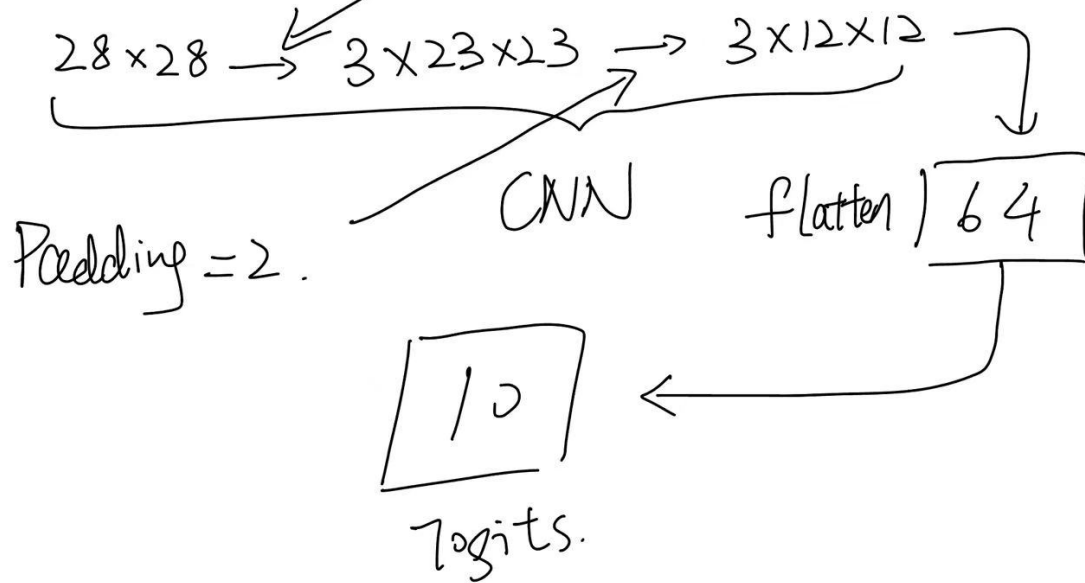
Gradients for Layer 3:
dW3:
[[-0.00011922 -0.0001192 -0.00011935 -0.00012044]
 [-0.00040692 -0.00040683 -0.00040736 -0.00041107]
 [ 0.00070526  0.00070511  0.00070603  0.00071246]]
db3:
[[-0.00023963]
 [-0.00081787]
 [ 0.00141751]]

Gradients for Layer 4:
dW4:
[[-0.10031159 -0.10184285 -0.10147129]
 [-0.06225442 -0.06320474 -0.06297415]]
db4:
[[-0.2021219 ]
 [-0.12543898]]

```

Task 2

B): CNN Using a 3×3 kernel, 3 chan output



```
In [ ]: import numpy as np
```

```
def relu(x):  
    return np.maximum(0, x)
```

```
def relu_deriv(x):  
    return (x > 0).astype(float)
```

```
class CNN:
```

```
    def __init__(self, kernel_size=3, num_filters=8, pool_size=2, num_classes=10):  
        """
```

```
        Initializes a simple CNN for MNIST.
```

```
        - Convolutional layer:  $3 \times 3$  filters, 'num_filters' filters.
```

```
        - Pooling layer:  $2 \times 2$  max pooling.
```

```
        - Fully connected layer: flattens and outputs 'num_classes' scores.
```

```
        The input image is assumed to be  $28 \times 28 \times 1$ .
```

```
        """
```

```
        self.kernel_size = kernel_size
```

```
        self.num_filters = num_filters
```

```
        self.pool_size = pool_size
```

```
        self.num_classes = num_classes
```

```
        # Initialize convolutional layer parameters.
```

```
        # Weight shape: (kernel_size, kernel_size, input_channels, num_filters)
```

```
        # For MNIST, input_channels = 1.
```

```
        self.W_conv = np.random.randn(kernel_size, kernel_size, 1, num_filters) * 0.01
```

```
        self.b_conv = np.zeros((1, 1, 1, num_filters))
```

```
        # After convolution:
```

```
        # Output height/width =  $28 - 3 + 1 = 26$ .
```

```
        # After  $2 \times 2$  pooling with stride 2:
```

```
        # Output height/width =  $26 // 2 = 13$ .
```

```
        self.conv_out_height = 26
```

```

self.conv_out_width = 26
self.pool_out_height = self.conv_out_height // pool_size
self.pool_out_width = self.conv_out_width // pool_size

# Fully connected layer parameters.
self.fc_input_dim = self.pool_out_height * self.pool_out_width * num_filters #
self.W_fc = np.random.randn(num_classes, self.fc_input_dim) * 0.01
self.b_fc = np.zeros((num_classes, 1))

def forward(self, X):
    """
    Forward propagation for the CNN.

    Arguments:
    - X: Input image, a numpy array of shape (28, 28, 1).

    Returns:
    - out_fc: Output of the fully connected layer (before any softmax),
              shape (num_classes, 1).
    Also stores intermediate values needed for backpropagation.
    """
    # Store input for backpropagation.
    self.X = X # shape: (28,28,1)

    # ---- Convolutional Layer Forward ----
    H, W, C = X.shape # (28, 28, 1)
    out_conv = np.zeros((H - self.kernel_size + 1, W - self.kernel_size + 1, self.num_filters))
    # Naive convolution (no padding, stride=1)
    for i in range(out_conv.shape[0]):
        for j in range(out_conv.shape[1]):
            # Extract the current 3x3 patch from input.
            patch = X[i:i+self.kernel_size, j:j+self.kernel_size, :] # shape: (3,3,1)
            for f in range(self.num_filters):
                conv_sum = np.sum(patch * self.W_conv[:, :, :, f]) + self.b_conv[0, f]
                out_conv[i, j, f] = conv_sum
    # Apply ReLU activation
    self.conv_out = relu(out_conv) # shape: (26,26,num_filters)

    # ---- Pooling Layer Forward (Max Pooling) ----
    H_conv, W_conv, F = self.conv_out.shape # (26,26,num_filters)
    H_pool = H_conv // self.pool_size # 13
    W_pool = W_conv // self.pool_size # 13
    out_pool = np.zeros((H_pool, W_pool, F))
    # To backpropagate through max pooling, record the mask.
    self.pool_mask = np.zeros_like(self.conv_out)
    for f in range(F):
        for i in range(H_pool):
            for j in range(W_pool):
                h_start = i * self.pool_size
                h_end = h_start + self.pool_size
                w_start = j * self.pool_size
                w_end = w_start + self.pool_size
                patch = self.conv_out[h_start:h_end, w_start:w_end, f]
                max_val = np.max(patch)
                out_pool[i, j, f] = max_val
                mask = (patch == max_val)
                self.pool_mask[h_start:h_end, w_start:w_end, f] = mask
    self.pool_out = out_pool

    self.flat = self.pool_out.reshape(-1, 1)
    Z_fc = np.dot(self.W_fc, self.flat) + self.b_fc
    self.fc_out = relu(Z_fc)
    return self.fc_out

def backward(self, dZ_fc, learning_rate=0.01):
    """

```

Backward propagation for the CNN.

Arguments:

- dZ_fc: Gradient of the loss with respect to the FC layer output (shape: (num_c
- learning_rate: Learning rate for parameter updates.

Returns:

- Gradients for parameters (for illustration).

```
"""
dW_fc = np.dot(dZ_fc, self.flat.T)
db_fc = dZ_fc
dflat = np.dot(self.W_fc.T, dZ_fc)
dpool = dflat.reshape(self.pool_out.shape)

dconv_out = np.zeros_like(self.conv_out) # shape: (26,26,num_filters)
H_pool, W_pool, F = self.pool_out.shape
for f in range(F):
    for i in range(H_pool):
        for j in range(W_pool):
            h_start = i * self.pool_size
            h_end = h_start + self.pool_size
            w_start = j * self.pool_size
            w_end = w_start + self.pool_size
            dconv_out[h_start:h_end, w_start:w_end, f] += self.pool_mask[h_start

dconv_out[self.conv_out <= 0] = 0

dW_conv = np.zeros_like(self.W_conv)
db_conv = np.zeros_like(self.b_conv)
H, W, C = self.X.shape
H_conv, W_conv, F = self.conv_out.shape # (26,26,num_filters)
for f in range(F):
    for i in range(H_conv):
        for j in range(W_conv):
            patch = self.X[i:i+self.kernel_size, j:j+self.kernel_size, :] # sha
            dW_conv[:, :, :, f] += patch * dconv_out[i, j, f]
            db_conv[0, 0, 0, f] += dconv_out[i, j, f]

# ---- Parameter Update ----
self.W_fc -= learning_rate * dW_fc
self.b_fc -= learning_rate * db_fc
self.W_conv -= learning_rate * dW_conv
self.b_conv -= learning_rate * db_conv

return dW_fc, db_fc, dW_conv, db_conv
```

```
np.random.seed(42)
cnn = CNN(kernel_size=3, num_filters=8, pool_size=2, num_classes=10)
```

```
X = np.random.randn(28, 28, 1)
out = cnn.forward(X)
print("CNN forward output (FC layer activations):\n", out)
```

```
dZ_fc = np.random.randn(10, 1)
grads = cnn.backward(dZ_fc, learning_rate=0.01)
```

```
print("\nParameter gradients computed during backward propagation:")
print("dW_fc shape:", grads[0].shape)
print("dW_conv shape:", grads[2].shape)
```

Task 3

```
In [4]: import random
import matplotlib.pyplot as plt
import kagglehub
import numpy as np
import struct
from array import array
from os.path import join
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
from torchvision import transforms
from tqdm import tqdm # Standard tqdm package
import numpy as np
from matplotlib import pyplot as plt
from sklearn.metrics import confusion_matrix
```

```
In [ ]: #####
## Install the training material
#####
path = kagglehub.dataset_download("hojjatk/mnist-dataset")
print("Path to dataset files:", path)
data_set_location = r"C:\Users\Eric\.cache\kagglehub\datasets\hojjatk\mnist-dataset\vers

Downloading from https://www.kaggle.com/api/v1/datasets/download/hojjatk/mnist-dataset?dataset_version_number=1...
100%|██████████| 22.0M/22.0M [00:00<00:00, 65.0MB/s]
Extracting files...

Path to dataset files: C:\Users\Eric\.cache\kagglehub\datasets\hojjatk\mnist-dataset\versions\1
```

```
In [5]: #####
## Define the loading function for the training material, obtained from the example give
#####
class MnistDataloader(object):
    def __init__(self, training_images_filepath, training_labels_filepath,
                 test_images_filepath, test_labels_filepath):
        self.training_images_filepath = training_images_filepath
        self.training_labels_filepath = training_labels_filepath
        self.test_images_filepath = test_images_filepath
        self.test_labels_filepath = test_labels_filepath

    def read_images_labels(self, images_filepath, labels_filepath):
        labels = []
        with open(labels_filepath, 'rb') as file:
            magic, size = struct.unpack(">II", file.read(8))
            if magic != 2049:
                raise ValueError('Magic number mismatch, expected 2049, got {}'.format(magic))
            labels = array("B", file.read())

        with open(images_filepath, 'rb') as file:
            magic, size, rows, cols = struct.unpack(">IIII", file.read(16))
            if magic != 2051:
                raise ValueError('Magic number mismatch, expected 2051, got {}'.format(magic))
            image_data = array("B", file.read())
        images = []
        for i in range(size):
            images.append([0] * rows * cols)
        for i in range(size):
            img = np.array(image_data[i * rows * cols:(i + 1) * rows * cols])
            img = img.reshape(28, 28)
            images[i][:] = img

        return images, labels
```

```

def load_data(self):
    x_train, y_train = self.read_images_labels(self.training_images_filepath, self.t
    x_test, y_test = self.read_images_labels(self.test_images_filepath, self.test_la
    return (x_train, y_train), (x_test, y_test)

#####
## Example loading
#####
input_path = r"C:\Users\Eric\.cache\kagglehub\datasets\hojatk\mnist-dataset\versions\1"
training_images_filepath = join(input_path, 'train-images-idx3-ubyte/train-images-idx3-u
training_labels_filepath = join(input_path, 'train-labels-idx1-ubyte/train-labels-idx1-u
test_images_filepath = join(input_path, 't10k-images-idx3-ubyte/t10k-images-idx3-ubyte')
test_labels_filepath = join(input_path, 't10k-labels-idx1-ubyte/t10k-labels-idx1-ubyte')

def show_images(images, title_texts):
    cols = 5
    rows = int(len(images)/cols) + 1
    plt.figure(figsize=(30,20))
    index = 1
    for x in zip(images, title_texts):
        image = x[0]
        title_text = x[1]
        plt.subplot(rows, cols, index)
        plt.imshow(image, cmap=plt.cm.gray)
        if (title_text != ''):
            plt.title(title_text, fontsize = 15);
        index += 1


mnist_dataloader = MnistDataloader(training_images_filepath, training_labels_filepath, t
(x_train, y_train), (x_test, y_test) = mnist_dataloader.load_data()

images_2_show = []
titles_2_show = []
for i in range(0, 10):
    r = random.randint(1, 60000)
    images_2_show.append(x_train[r])
    titles_2_show.append('training image [' + str(r) + '] = ' + str(y_train[r]))

for i in range(0, 5):
    r = random.randint(1, 10000)
    images_2_show.append(x_test[r])
    titles_2_show.append('test image [' + str(r) + '] = ' + str(y_test[r]))

show_images(images_2_show, titles_2_show)

```

```
In [ ]: # -----
# Set up file paths for the MNIST dataset (adjust these paths as needed)
# -----
input_path = r"C:\Users\Eric\.cache\kagglehub\datasets\hojjatk\mnist-dataset\versions\1"
training_images_filepath = join(input_path, 'train-images-idx3-ubyte/train-images-idx3-ubyte')
training_labels_filepath = join(input_path, 'train-labels-idx1-ubyte/train-labels-idx1-ubyte')
test_images_filepath = join(input_path, 't10k-images-idx3-ubyte/t10k-images-idx3-ubyte')
test_labels_filepath = join(input_path, 't10k-labels-idx1-ubyte/t10k-labels-idx1-ubyte')

# Load the MNIST data using your loader
mnist_dataloader = MnistDataloader(training_images_filepath, training_labels_filepath,
                                   test_images_filepath, test_labels_filepath)
(x_train, y_train), (x_test, y_test) = mnist_dataloader.load_data()
print("MNIST data loaded.")

# Convert the loaded images and labels to NumPy arrays
x_train = np.array(x_train) # shape: (num_train, 28, 28)
x_test = np.array(x_test) # shape: (num_test, 28, 28)
y_train = np.array(y_train) # shape: (num_train,)
y_test = np.array(y_test) # shape: (num_test,)

# -----
# Convert the data to PyTorch tensors (and add a channel dimension)
# -----
train_images_tensor = torch.tensor(x_train, dtype=torch.float32).unsqueeze(1) # shape:
train_labels_tensor = torch.tensor(y_train, dtype=torch.long)
test_images_tensor = torch.tensor(x_test, dtype=torch.float32).unsqueeze(1) # shape:
test_labels_tensor = torch.tensor(y_test, dtype=torch.long)

# Create a TensorDataset for training
train_dataset = TensorDataset(train_images_tensor, train_labels_tensor)

# -----
# Define the CNN Model (using PyTorch modules)
# -----
class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv = nn.Conv2d(in_channels=1, out_channels=3, kernel_size=6)
```

```

self.relu = nn.ReLU()
self.pool = nn.MaxPool2d(kernel_size=2, padding=1)

self.fc1 = nn.Linear(3 * 12 * 12, 60)
self.sigmoid = nn.Sigmoid()
self.fc2 = nn.Linear(60, 10)
self.softmax = nn.Softmax(dim=1)

def forward(self, x):
    x = self.conv(x)
    x = self.relu(x)
    x = self.pool(x)
    x = torch.flatten(x, 1)
    x = self.fc1(x)
    x = self.sigmoid(x)
    x = self.fc2(x)
    x = self.softmax(x)
    return x

def train_all(self, dataset, batch_size, lr, epochs):
    loader = DataLoader(dataset, batch_size=batch_size, shuffle=True)
    crit = nn.CrossEntropyLoss()
    opt = optim.Adam(self.parameters(), lr=lr)
    losses = []
    total_steps = len(loader) * epochs
    with tqdm(total=total_steps, desc="CNN Training") as pbar:
        for epoch in range(epochs):
            self.train()
            epoch_loss = 0.0
            for images, labels in loader:
                images, labels = images.to(device), labels.to(device)
                opt.zero_grad()
                outputs = self(images)
                loss = crit(outputs, labels)
                loss.backward()
                opt.step()
                epoch_loss += loss.item()
                pbar.update(1)
            losses.append(epoch_loss)
            #print(f"Epoch {epoch+1}/{epochs}: Loss = {epoch_loss:.4f}")
    return losses

# -----
# Define the MLP Model (using PyTorch modules)
# -----
class MLP(nn.Module):
    def __init__(self):
        super(MLP, self).__init__()
        self.fc1 = nn.Linear(28 * 28, 500)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(500, 300)
        self.sigmoid = nn.Sigmoid()
        self.fc3 = nn.Linear(300, 10)
        self.softmax = nn.Softmax(dim=1)

    def forward(self, x):
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = self.relu(x)
        x = self.fc2(x)
        x = self.sigmoid(x)
        x = self.fc3(x)
        x = self.softmax(x)
        return x

    def train_all(self, dataset, batch_size, lr, epochs):

```

```

        loader = DataLoader(dataset, batch_size=batch_size, shuffle=True)
        crit = nn.CrossEntropyLoss()
        opt = optim.Adam(self.parameters(), lr=lr)
        losses = []
        total_steps = len(loader) * epochs
        with tqdm(total=total_steps, desc="MLP Training") as pbar:
            for epoch in range(epochs):
                self.train()
                epoch_loss = 0.0
                for images, labels in loader:
                    images, labels = images.to(device), labels.to(device)
                    opt.zero_grad()
                    outputs = self(images)
                    loss = crit(outputs, labels)
                    loss.backward()
                    opt.step()
                    epoch_loss += loss.item()
                    pbar.update(1)
                losses.append(epoch_loss)
        return losses

def conf_matrix(model):
    model.eval()
    with torch.no_grad():
        outputs = model(test_images_tensor.to(device))
        preds = torch.argmax(outputs, dim=1).cpu().numpy()
    conf = confusion_matrix(y_test, preds)
    fig, ax = plt.subplots()
    im = ax.imshow(conf, cmap='Reds')
    ax.set_xticks(np.arange(conf.shape[1]))
    ax.set_yticks(np.arange(conf.shape[0]))
    ax.set_xticklabels(range(10))
    ax.set_yticklabels(range(10))
    plt.setp(ax.get_xticklabels(), rotation=45, ha="right", rotation_mode="anchor")

    thresh = conf.max() / 2.0
    for i in range(conf.shape[0]):
        for j in range(conf.shape[1]):
            color = "white" if conf[i, j] > thresh else "black"
            ax.text(j, i, f"{conf[i, j]:d}", ha="center", va="center", color=color)

    ax.set_xlabel('Predicted')
    ax.set_ylabel('Actual')
    plt.title("Confusion Matrix")
    plt.colorbar(im)
    plt.show()

```

MNIST data loaded.

```

In [50]: #####
# Initialize the CNN and MLP instance
#####
gpu = "cuda"
batch = 512
epoch = 100
iteration_step = 0.0001

device = torch.device(gpu)
cnn = CNN().to(device)

device = torch.device(gpu)
mlp = MLP().to(device)


```


```
print(mlp.fc1.weight.shape)
print(mlp.fc1.bias.shape)
```

```
torch.Size([500, 784])
torch.Size([500])
```

```
In [ ]: cnn_losses = cnn.train_all(train_dataset, batch_size=batch, lr=iteration_step, epochs=ep
plt.plot(cnn_losses)
plt.title("CNN Loss Curve")
plt.xlabel("Epoch")
plt.ylabel("Total Loss")
plt.show()
conf_matrix(cnn)
```



CNN Training: 100%|██████████| 11800/11800 [01:01<00:00, 192.43it/s]





```
In [ ]: mlp_losses = mlp.train_all(train_dataset, batch_size=batch, lr=iteration_step, epochs=ep
plt.plot(mlp_losses)
plt.title("MLP Loss Curve")
plt.xlabel("Epoch")
plt.ylabel("Total Loss")
plt.show()
conf_matrix(mlp)
```

MLP Training: 100%|██████████| 11800/11800 [00:56<00:00, 210.57it/s]



In []: